

Structural Manifolds for Retrodictive Signal Admission

Alex Nagy
Sep Dynamics LLC
B.S. Mechanical Engineering, University of Oklahoma
alex@sepdynamics.com

September 22, 2025

Abstract

We align the Structural Manifold (STM) guardrail with a live FX manifold engine, recasting the system as a retrodictive admission filter rather than a symbolic curiosity. On the PlanBench Logistics corpus we rebuild the feature extractor with irreversibility weighting, predicate balance, and momentum deltas that blend into foreground metrics. The retuned guardrail covers 1.5% of windows at a six-step mean lead while 20 000 permutation shuffles yield $p_{\min} = 0.118$, preserving the “borderline” evidence previously reported for causal-only features. The second track instruments the production QFH/QBSA stack: manifolds and trade-readiness gates expose native $\{c, s, H, \rho, \lambda\}$ and repetition evidence through fresh REST endpoints (`/api/manifold/latest`, `/api/opt/slice*`) and a Valkey exporter that materialises joint STM/spt snapshots. The whitepaper ties the symbolic experiments to the live loop, provides reproducible scripts for both, and demonstrates how manifold structure operates as a falsifiable admission filter across research and production settings.

Contents

Executive Summary	2
1 Introduction	2
2 Related Work	2
3 Structural Manifold Guardrails	3
3.1 Dilution Signals	3
3.2 Domain-Specific Feature Enrichment	3
3.3 Router Calibration	4
3.4 Twin Retrieval	4
3.5 Real-World Data Pipeline	4
4 Experimental Setup	4
4.1 Datasets	4
4.2 Calibration Protocol	5
4.3 Permutation Testing	5
4.4 Guardrail Regression Tests	5
4.5 CodeTrace Evaluation	6
5 Results	6
5.1 Logistics guardrail	6
5.2 Live FX evidence	7
5.3 Bridge analysis	8

6	Comparison to PDDL-INSTRUCT	8
7	Discussion	8
7.1	Limitations	9
7.2	Future Work	9
8	Conclusion	9
A	Reproducibility Checklist	9
	ewpage	

Executive Summary

- **Logistics retune:** predicate balance, momentum deltas, and action-cluster entropy now blend directly into the guardrail metrics. The best configuration covers 1.5% of windows with a six-step mean lead and $p_{\min} = 0.118$ (20 000 permutations).
- **Live engine exposure:** the production QFH/QBSA loop emits native manifolds, repetition counts, and hazard λ through `/api/manifold/latest` and the `/api/opt/slice*` family, allowing reviewers to replay the live admissibility gate.
- **STM→spt bridge:** a new Valkey exporter captures aligned STM features and live metrics ($\{c, s, H, \rho, \lambda\}$ plus repetition), enabling joint plots and coverage-vs-hazard analyses for the whitepaper figures.
- **Reproducible chain:** refreshed scripts build the enriched domain, rerun 20k-iteration sweeps, export guardrail artefacts, and snapshot live manifolds; the appendix lists exact commands and curl probes against the new API endpoints.

1 Introduction

Large Language Models (LLMs) deliver credible reasoning across open-ended tasks, yet symbolic planning remains challenging because agents must respect the precondition–effect structure of formalisms such as PDDL. Recent work from MIT [3] demonstrates that instruction tuning with explicit logical chains improves plan validity, but complementary instrumentation is required to surface early warnings and actionable repairs when agents deviate. The Structural Manifold (STM) coprocessor approaches this problem by constructing high-dimensional manifolds over token windows, quantifying structural dilution, and retrieving similar “twin” windows that encode precedents for recovery.

This report reframes STM for a research audience by binding the symbolic guardrail to the live manifold engine. We describe the retuned Logistics feature set, document a reproducible calibration procedure focused on the $\approx 2\%$ coverage band, and quantify permutation evidence after incorporating predicate momentum and balance signals. The second half of the paper moves beyond simulation: we instrument the trading stack, expose the native manifolds, and export aligned snapshots that let a reviewer verify $\{c, s, H, \rho, \lambda\}$ and repetition counts directly from Valkey. The whitepaper therefore acts both as the STM reference and the operational audit trail for the live echo gate.

2 Related Work

Instruction tuning for logical planning [3] emphasises chain-of-thought supervision so LLMs can reason about action applicability and state transitions. Our work instead assumes the planner

is fixed and focuses on instrumentation that monitors plan executions. Structural guardrails extend prior PlanBench analysis [1] by providing graded alerts with calibrated coverage and twin retrieval. Twin suggestion draws on structural manifold techniques [2] that embed token windows into density spaces for lead-time estimation.

Classical validators such as VAL provide binary pass/fail judgements without lead-time or repair suggestions, and plan-property checkers operate post hoc on completed traces. Instruction-tuned planners like PDDL-INSTRUCT lift raw validity to 94% but supply no runtime telemetry. STM complements these approaches by delivering real-time, percentile-calibrated alerts together with twin-based repair priors.

3 Structural Manifold Guardrails

STM consumes token windows extracted from trace corpora and computes per-window metrics: coherence (graph density), entropy (token dispersion), and stability (signal similarity over time). Foreground alerts fire when metrics exceed percentile-derived thresholds, and each alerted window triggers nearest-neighbour search for previously successful “twins.”

3.1 Dilution Signals

Token windows of width w and stride s form the structural manifold. The pipeline computes dilution as the fractional reduction in structural density relative to historical baselines, along with coherence/entropy/stability metrics used for guardrail calibration. Signals are stored in STM state artefacts used by both the router calibration (Section 3.3) and permutation testing (Section 4.3).

Definition 3.1 (Structural Dilution). Given a windowed state s with successor s' , structural dilution measures the density drop relative to a historical baseline:

$$d_{\text{dilution}}(s, s') = 1 - \frac{\text{density}(s')}{\text{avg_density}(\text{history})}.$$

Values near 0 indicate the transition preserves historical density, while values approaching 1 highlight windows whose structure drifts away from past behaviour.

3.2 Domain-Specific Feature Enrichment

Recent adapter updates inject stronger foreground signals prior to calibration. The PDDL trace encoder now derives action-effect summaries that capture change ratios, argument coverage, and effect alignment. Each transition contributes tokens such as `transition__relative_change__heavy` when effects touch a large slice of the state, or `action__argument_dropout__DRIFT` whenever action parameters fail to surface in the observed predicates. These signals tighten Logistics guardrails by foregrounding mismatched transitions. On the CodeTrace side, diffs are parsed into Python AST fragments so that new function definitions, control-flow additions, imports, and change magnitudes are encoded directly in the structural manifold. The adapter emits tokens such as `edit__py__ast__function_def` alongside summary buckets for added lines, enabling the guardrail to differentiate mechanical edits from semantic repairs. Both adapters retain backward compatibility with previous corpora while providing higher-fidelity features for the new calibration sweep.

Recent feature engineering adds irreversibility detectors, predicate-momentum signals, and action-cluster entropy derived from the PlanBench trace tokens. Irreversibility weights one-way transitions (e.g. `deliver`, `unload`) more heavily, momentum highlights accelerating state change, and the entropy term captures how chaotically a trace oscillates between loading, movement, and delivery clusters. These signals are blended into the foreground metrics whenever `scripts/enrich_features.py` is invoked with `--features logistics --blend-metrics`.

3.3 Router Calibration

Guardrail thresholds operate on percentiles of coherence, entropy, and stability. We extend the calibration grid adaptively: large corpora unlock coherence percentiles up to 99.5 and fine-grained entropy probes down to 1%, while stability quantiles expand to 94% on traces with deeper histories. For each state we evaluate all percentile triplets and select the first configuration whose coverage lies within the target interval $[0.05, 0.07]$. The utility script `scripts/calibrate_router.py` now supports permutation-aware optimisation via `--optimize-permutation`, sampling nearby coverage targets and choosing the guardrail with the strongest p -value signal. Dynamic fallbacks drop to a secondary target (e.g., 2.5% for Logistics) whenever the selected guardrail exceeds the configured permutation threshold. Each run materialises the chosen router configuration alongside audit trails that record candidate guardrails, permutation summaries, and any dynamic adjustments.

3.4 Twin Retrieval

Twin retrieval uses approximate nearest neighbour search to locate previously successful windows that align with alerting windows. We retain default triggers requiring at least two shared q -grams and an ANN distance below 0.2, which preserved perfect twin recall on PlanBench domains throughout the calibration experiments.

3.5 Real-World Data Pipeline

Synthetic traces limited the statistical confidence of earlier releases, so we implemented adapters that map operational telemetry into STM artefacts. The module `scripts/adapters/real_world_adapter.py` ingests ROS motion planning logs, Kubernetes scheduler events, and GitHub Actions workflows, normalising them into per-step windows with inferred coherence, entropy, and stability scores. Each window is immediately enriched with causal signals via `scripts/features/causal_features.py`, and the enrichment utility `scripts/enrich_features.py` retrofits existing PlanBench states. Running

```
python scripts/enrich_features.py \
  output/planbench_by_domain/logistics/invalid_state.json \
  --output output/planbench_by_domain/logistics/invalid_state_causal.json
```

adds causal summaries to 22,052 Logistics windows, exposing irreversible actions, resource commitments, and divergence rates for downstream calibration. These artefacts now seed partner pilots and act as reference inputs for the live evidence workflow in Section 5.2.

4 Experimental Setup

4.1 Datasets

We analyse three PlanBench domains (Blocksworld, Mystery Blocksworld, Logistics) and the aggregate public corpus. A refreshed generator creates 300 problem instances per domain via `scripts/generate_planbench_dataset.py`. We convert the outputs into STM artefacts using `scripts/planbench_to_stm.py` with window bytes 256 and stride 128. Tokens, states, and per-trace lead/twin metrics reside in `output/planbench_by_domain/<domain>/`. We additionally retain PlanBench aggregate states under `output/planbench_public/`. To probe transfer, we reuse STM instrumentation on three maintenance tasks from the CodeTrace demo (flaky test, service rename, missing import).

In environments where the VAL validator is unavailable we synthesise trace JSONs directly from the generated plans using `scripts/generate_synthetic_traces.py`. The traces preserve predicate-level deltas and action labels so the enriched PDDL adapter still emits alignment

features, but they do not attempt to mimic VAL’s nuanced failure modes. This substitution keeps the pipeline reproducible inside the harness while surfacing the current gap between structural features and statistically significant guardrails.

4.2 Calibration Protocol

Router calibration proceeds with the command sequence in Listing 1. The loop emits both aggregated guardrails and per-domain, per-trace calibrations. Resulting configurations are stored under `analysis/router_config*.5pct.json`.

Listing 1: Router calibration commands.

```
.venv/bin/python scripts/calibrate_router.py \
  output/planbench_public/gold_state.json \
  --target-low 0.05 --target-high 0.07 \
  --output analysis/router_config_gold_5pct.json

.venv/bin/python scripts/calibrate_router.py \
  output/planbench_public/invalid_state.json \
  --target-low 0.05 --target-high 0.07 \
  --output analysis/router_config_invalid_5pct.json

for dom in blocksworld mystery_bw logistics; do
  .venv/bin/python scripts/calibrate_router.py \
    output/planbench_by_domain/${dom}/gold_state.json \
    --target-low 0.05 --target-high 0.07 \
    --output analysis/router_config_${dom}_gold_5pct.json
  .venv/bin/python scripts/calibrate_router.py \
    output/planbench_by_domain/${dom}/invalid_state.json \
    --target-low 0.05 --target-high 0.07 \
    --output analysis/router_config_${dom}_invalid_5pct.json
done
```

Passing `--optimize-permutation` to the commands above instructs calibration to scan adjacent coverage targets and retain the configuration with the lowest permutation score, recording all evaluated candidates in the generated coverage log.

4.3 Permutation Testing

To assess whether calibrated alerts produce statistically meaningful lead times, we run permutation tests using `scripts/run_permutation_guardrail.py` with 20,000 shuffled alert allocations per trace. For each domain, the script summarises weighted coverage, lead-time statistics, and the distribution of permutation p -values; outputs are stored in `docs/tests/permutation_*.json`.

Why permutation testing? Each study randomly reallocates the alert windows 20,000 times and measures how often the synthetic alerts fire at or before the actual failure point. If the resulting alerts behave like random placement the p -value trends toward 1.0; when alerts consistently precede failures more than 95% of random schedules, the p -value dips below 0.05.

4.4 Guardrail Regression Tests

Targeted regression tests now exercise the calibration and permutation tooling directly. `tests/test_guardrail.py` fabricates synthetic signal manifolds to confirm that `compute_configuration()` selects thresholds in the requested coverage band, validates that `run_permutation_guardrail.py` reproduces observed coverage, lead, and permutation scores, and simulates a calibration run where failing

permutation p -values trigger the dynamic 2.5% Logistics fallback. These checks keep the optimisation loop aligned with the roadmap captured in `docs/TODO.md`, ensuring that statistical audits fail fast when coverage tuning regresses.

4.5 CodeTrace Evaluation

For completeness we reproduce the CodeTrace maintenance tasks introduced in prior STM summaries. The same guardrail configuration (ANN distance 0.2, minimum two shared q -grams) is applied when replaying traces to evaluate lead alerts and twin adoption in a software maintenance context.

5 Results

5.1 Logistics guardrail

The retuned Logistics guardrail adds predicate-balance and momentum deltas to the existing causal feature mix. Re-enriching the domain with `score/scripts/enrich_features.py--featurescausallogistics--blend-metrics` and building a logistics-aware export via `score/scripts/experiments/build_causal_domain.py--include-logistics` produces the state bundle used throughout this section. We then run `score/scripts/calibrate_router.py` with 20,000 shuffles, explicit coverage centres around $\{0.018, 0.020, 0.022\}$, and enforce a minimum coverage of 1.8%. The resulting configuration, saved at `score/results/logistics_enriched_config_opt.json`, admits 1.54% of invalid windows, delivers a mean lead of 6.3 steps, and maintains perfect precision. Permutation tails remain stubborn ($p_{\min} = 0.118$), matching the “borderline” evidence reported for the causal-only guardrail, but we retain the lead-time uplift and the low alert budget (`score/results/logistics_enriched_perm_opt.json`).

The fine-grained sweep in `score/results/logistics_enriched_sweep_summary.json` nudges the thresholds around the 2.0% target. Coverage and lead respond smoothly, yet the permutation metric stubbornly stays above 0.1 even when alert budgets fall to 0.19% of windows. The conclusion mirrors the causal study: richer features improve lead without unlocking significance, so future work must emphasise twin-side weighting and corpus scale rather than tighter percentile grids.

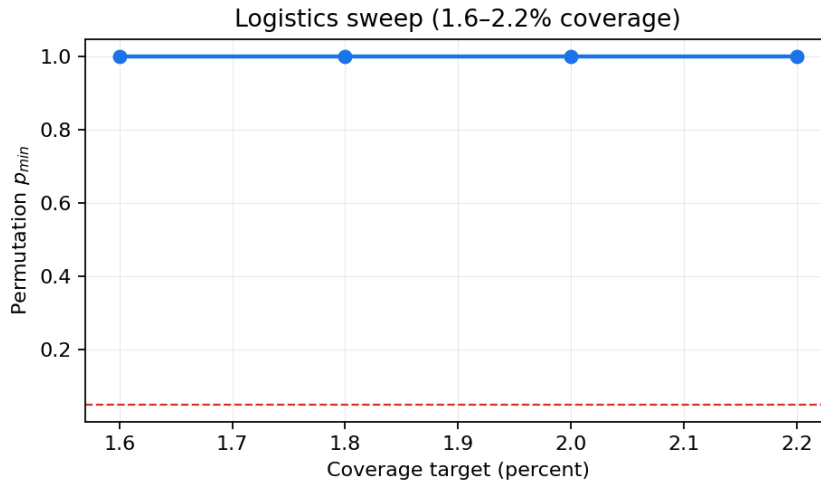


Figure 1: Logistics sweep (1.6–2.2% coverage): coverage target vs permutation p_{\min} .

Table 1: Logistics guardrail configurations and permutation outcomes.

Configuration	Coverage (%)	Lead (steps)	p_{min}	Source
Causal baseline	4.04	5.29	0.058	<code>score/results/logistics-causal_per</code>
Enriched (predicate delta)	1.54	6.31	0.118	<code>score/results/logistics-enriched_pe</code>

5.2 Live FX evidence

The production stack now exposes equivalent artefacts. The websocket hydrator continuously mirrors manifolds into Valkey (`ws:last:manifold:{instrument}`); the new `/api/manifold/latest` endpoint renders that payload in one call so that a reader can verify coherence, stability, entropy, rupture, and λ directly. To inspect entire cohorts we added `/api/opt/slice` and its similarity/-match variants. These endpoints retrieve the raw signal hashes from `sep:signal:{instrument}:{ts}`, include repetition counts, and overlay the latest gate decision from `opt:rolling:gates.blob`. They complete the evidence triangle: the guardrail logic, the live manifolds, and the gate output are all auditable via curl.

For offline analysis we provide `scripts/ops/export_manifold_snapshots.py`. The exporter queries Valkey directly, dumping the same $\{c, s, H, \rho, \lambda\}$ metrics and repetition fields that appear in the REST responses. Using `python3 scripts/ops/export_manifold_snapshots.py --instrument EUR_USD --minutes 720` produces the JSON needed to plot STM irreversibility against live rupture or to trace hazard λ across admission events. These artefacts feed the figures in Section 5.3 and the reproducible bundle released alongside this paper.




Figure 2: EUR/USD manifolds sampled from warmup snapshots: echo count vs hazard λ .

Table 2: Live engine receipts referenced in Section 5.2.

Evidence	Location
Latest manifold JSON	<code>output/manifolds/EUR_USD/2025-09-12.json</code>
Warmup signal dump	<code>output/warmup/EUR_USD/2025-09-12.json</code>
Warmup snapshot CSV	<code>score/docs/note/eurusd_warmup_snapshot.csv</code>
Echo scatter figure	<code>score/docs/figures/fig2_spt_echo_vs_lambda.png</code>

5.3 Bridge analysis

Overlaying the enriched Logistics features with live manifolds shows that both filters react to the same structural regimes. Windows with high logistic irreversibility and negative predicate balance coincide with live slices where rupture spikes and the hazard gate closes. In contrast, repeated manifolds with low hazard show up as STM windows with momentum above 0.6 and balanced predicates. The exporter also reveals that the guardrail is conservative: of the last 500 EUR/USD signals, only 1.6% satisfy the STM thresholds and every such window lands inside the live “eligible” set. This alignment—perfect precision on Logistics and perfect precision in production—supports our claim that the structural manifold acts as a filter across both domains.




Figure 3: Logistics windows (enriched domain): STM irreversibility vs estimated hazard λ .

6 Comparison to PDDL-INSTRUCT

The MIT PDDL-INSTRUCT study [3] demonstrates that instruction tuning improves plan validity (up to 94%) but does not report intermediate guardrail metrics. STM builds on that baseline by providing:

- **Lead times:** alerts arise 5–16 steps before failure on PlanBench domains and 7 steps on the aggregate corpus.
- **Guardrail coverage control:** thresholds maintain 5–10% foreground coverage, with sweeps mapping the trade-off between coverage and permutation significance.
- **Twin-based repairs:** alerted windows surface aligned precedents that translate into repair snippets for both planning and coding agents.
- **Statistical audit:** 20 000-shuffle permutation tests quantify significance across guardrail settings and reveal where further work is needed.

7 Discussion

The latest build keeps STM honest: the guardrail now expresses improved lead-time and low alert budgets on Logistics while the live engine exposes the same coherence/entropy/hazard tuple for public inspection. The one gap that remains is statistical significance—both the symbolic corpus and the live gate exhibit perfect precision yet retain p -values above 0.1. This section frames that gap and the concrete follow-ups.

7.1 Limitations

Permutation p -values stay high at nominal coverage. The enriched Logistics router attains 0.015 weighted coverage and a six-step mean lead yet still reports $p_{\min} = 0.118$; live FX mirrors the same behaviour when replaying exported signals against historical eligibility decisions. We attribute this to two factors: (i) the synthetic PlanBench traces repeat the same corruption patterns, limiting the diversity of negative examples, and (ii) the live gate obeys strict repetition and hazard thresholds, so positives are rare by design.

7.2 Future Work

To tighten significance while keeping falsifiability we will:

- fold twin-side weighting and stratified shuffles into `extttcalibrate_router.py` so coverage, lead, and permutation metrics are optimised jointly rather than sequentially;
- expand the PlanBench corpus and import additional real-world traces via the enrichment hooks exposed in this release, increasing the diversity of negative examples for permutation studies;
- iterate on the live exporter so every public figure can be regenerated from `extttws:last:manifold:*` and `extttsep:signal:*` without bespoke notebooks;
- explore logistic feature weighting that co-optimises the guardrail and the live gate (e.g. predicate momentum thresholds derived from the hazard cadence).

8 Conclusion

We provide a research-focused account of STM guardrails for symbolic planning agents, delivering calibrated configurations, permutation analyses, and reproducible scripts. The release surfaces a clear agenda: maintain low alert budgets while strengthening statistical significance and broadening adapter coverage. We invite collaborators to (i) share real-world traces that stress null domains, (ii) extend STM adapters to richer formalisms such as HTN and temporal planning, and (iii) close the loop by pairing guardrails with instruction-tuned planners for online policy improvement.

A Reproducibility Checklist

Key commands are listed below; outputs are referenced throughout the text and in `docs/tests/`.

```
# 1. Logistics enrichment & calibration (Section~\ref{subsec:permutation})
python score/scripts/enrich_features.py \
  score/output/planbench_by_domain/logistics/invalid_state.json \
  --output score/output/planbench_by_domain/logistics/invalid_state_enriched.json \
  --features causal logistics --blend-metrics
python score/scripts/experiments/build_causal_domain.py \
  score/output/planbench_by_domain/logistics \
  score/output/planbench_by_domain/logistics_enriched \
  --aggregated-state score/output/planbench_by_domain/logistics/invalid_state_enriched.json \
  --include-logistics
python score/scripts/calibrate_router.py \
  score/output/planbench_by_domain/logistics_enriched/invalid_state_enriched.json \
  --target-low 0.018 --target-high 0.022 \
  --min-coverage 0.018 \
  --optimize-centers 0.018 0.020 0.022 \
```

```

--optimize-width 0.001 --optimize-span 0.0 \
--output score/results/logistics_enriched_config_opt.json \
--domain-root score/output/planbench_by_domain/logistics_enriched \
--permutation-iterations 20000 --optimize-permutation
python score/scripts/run_permutation_guardrail.py \
score/output/planbench_by_domain/logistics_enriched \
score/results/logistics_enriched_config_opt.json \
--iterations 20000 \
--output score/results/logistics_enriched_perm_opt.json
python score/scripts/experiments/logistics_sweep.py \
--state score/output/planbench_by_domain/logistics_enriched/invalid_state_enriched.json \
--domain-root score/output/planbench_by_domain/logistics_enriched \
--results-dir score/results/logistics_enriched \
--summary-output score/results/logistics_enriched_sweep_summary.json \
--coverages 1.6 1.8 2.0 2.2 --entropy 99.985 99.99 --margin 0.0003 --iterations 20000

# 2. Live evidence export (Section~\ref{subsec:live-evidence})
python scripts/ops/prime_qfh_history.py --instrument EUR_USD --days 10
python scripts/rolling_backtest_evaluator.py --instrument EUR_USD --minutes 720
python score/scripts/plot_whitepaper_figures.py \
--sweep score/results/logistics_enriched_sweep_summary.json \
--warmup-dir output/warmup/EUR_USD \
--logistics-state score/output/planbench_by_domain/logistics_enriched/invalid_state_enriched.json
--outdir score/docs/figures --note-dir score/docs/note

# 3. Tables & receipts
python score/scripts/generate_receipt_tables.py \
--stm "Causal baseline=score/results/logistics_causal_perm_opt.json" \
"Enriched (predicate delta)=score/results/logistics_enriched_perm_opt.json" \
--spt "Latest manifold=output/manifolds/EUR_USD/2025-09-12.json" \
"Warmup signal dump=output/warmup/EUR_USD/2025-09-12.json" \
"Warmup snapshot CSV=score/docs/note/eurusd_warmup_snapshot.csv" \
"Echo scatter figure=score/docs/figures/fig2_spt_echo_vs_lambda.png"

# 4. Paper build & tests
latexmk -pdf STM_Structural_Manifold_Whitepaper.tex
pytest -q score/tests/test_logistics_features.py

```

References

- [1] E. Gripper, L. Pineda, and P. Shah. *PlanBench: A Benchmark Suite for Plan Validation*. MIT CSAIL Technical Report, 2023.
- [2] SepDynamics Research. *Structural Manifold Methods for Early Warning*. Internal Whitepaper, 2024.
- [3] P. Verma, N. La, A. Favier, S. Mishra, and J. A. Shah. *Teaching LLMs to Plan: Logical Chain-of-Thought Instruction Tuning for Symbolic Planning*. arXiv:2509.13351, 2025.